

Case-Study Day 11.2

Task 1:

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)
2. Create a new table called TRANSACTIONS_COUNT. This table should have 3 fields - custid, fname and count. (Again, to be done in module 8)
3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).
4. Now lets make the TRANSACTIONS_COUNT table Hbase complaint. In the sence, use Ser Des And Storate handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)
5. Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically (This has to be done in module 10)
6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

Answers:

hive -S

(Starting HIVE in suppress mode to avoid excess info)

create database acadgilddb;

(Creating a database by the name acadgilddb)

show databases;

(Listing the databases present)

acadgilddb

use acadgilddb;

(Use acadgilddb to create both the tables **CUSTOMER & TRANSACTIONS.**)

ScreenShot:

```
hive> use acadgilddb;
hive> show tables
> ;
hive>
hive> CREATE TABLE CUSTOMER(
> custid INT,
> fname STRING,
> lname STRING,
> age INT,
> profession STRING)
> row format delimited fields terminated by ',';
hive> LOAD DATA LOCAL INPATH '/home/acadgild/Desktop/TestHadoop/hive/custs.txt'
> into table CUSTOMER;
hive> CREATE TABLE TRANSACTIONS (
> txnno INT,
> txndate STRING,
> custno INT,
> amount DOUBLE,
> category STRING,
> product STRING,
> city STRING,
> state STRING,
> spendby STRING)
> row format delimited fields terminated by ',';
hive> LOAD DATA LOCAL INPATH '/home/acadgild/Desktop/TestHadoop/hive/txn.txt'
> into table TRANSACTIONS;

hive> select * from CUSTOMER;
101 Amitabh Bacchan 65 Actor
102 Sharukh Khan 45 Doctor
103 Akshay Kumar 38 Dentist
104 Anubahv kumar 58 Business
105 Pawan Trivedi 34 service
106 Aamir Null 42 scientest
107 Salman Khan 43 Surgen
108 Ranbir Kapoor 26 Industrialist
hive> select * from TRANSACTIONS;
97834 05/02/2018 101 965.0 Entertainment Movie Pune Maharashtra Daughter
98396 12/01/2018 102 239.0 Food Grocery Patna Bihar Self
34908 06/01/2018 101 875.0 Travel Air Bangalore Karnataka Spouse
70958 17/02/2018 104 439.0 Food Restaurant Delhi Delhi Wife
9874 21/01/2018 105 509.0 Entertainment Park Kolkata West Bengal NULL
94585 19/01/2018 106 629.0 Rent House Hyderabad Telangana Self
45509 20/01/2018 107 953.0 Travel Rail Chennai Tamil Nadu Brother
7864 01/02/2018 108 569.0 Rent Parking Goa Goa Wife
```

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)

Ans:

select t.custno,c.fname,count(txnno) **from** TRANSACTIONS t **join** CUSTOMER c on t.custno=c.custid **group by** t.custno,c.fname;

(listing out names of all such customers who have done a transaction by joining both the tables on cust id).

ScreenShot:

```
hive> select c.fname, t.custno, count(txnno) from TRANSACTIONS t join CUSTOMER c on t.custno=c.custid group by t.custno,c.fname;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution
ark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180806220013_a0f977ee-f25f-4586-82a1-d49a9fea6a8c
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/im
nder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/or
icLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-08-06 22:00:29 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-06 22:00:33 Dump the side-table for tag: 1 with group count: 8 into file: file:/tmp/acadgild/4f191ae7-e14d-4cc2-b22b
e_2018-08-06_22-00-13_413_1094874905996884690-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile01--.hashtable
2018-08-06 22:00:33 Uploaded 1 File to: file:/tmp/acadgild/4f191ae7-e14d-4cc2-b22b-433fe808fe5b/hive_2018-08-06_22-00-13_413
690-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile01--.hashtable (469 bytes)
2018-08-06 22:00:33 End of local task; Time Taken: 3.492 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533488475608_0003, Tracking URL = http://localhost:8088/proxy/application_1533488475608_0003/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533488475608_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-06 22:00:52,746 Stage-2 map = 0%, reduce = 0%
2018-08-06 22:01:05,935 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 3.83 sec
2018-08-06 22:01:19,987 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 6.81 sec
MapReduce Total cumulative CPU time: 6 seconds 810 msec
Ended Job = job_1533488475608_0003
MapReduce Jobs Launched:

2018-08-06 22:00:29 Starting to launch local task to process map join; maximum memory = 5189795
2018-08-06 22:00:33 Dump the side-table for tag: 1 with group count: 8 into file: file:/tmp/acadgild
e_2018-08-06_22-00-13_413_1094874905996884690-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile01--.hasht
2018-08-06 22:00:33 Uploaded 1 File to: file:/tmp/acadgild/4f191ae7-e14d-4cc2-b22b-433fe808fe5b/hive
690-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile01--.hashtable (469 bytes)
2018-08-06 22:00:33 End of local task; Time Taken: 3.492 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533488475608_0003, Tracking URL = http://localhost:8088/proxy/application_1533488475
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533488475608_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-06 22:00:52,746 Stage-2 map = 0%, reduce = 0%
2018-08-06 22:01:05,935 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 3.83 sec
2018-08-06 22:01:19,987 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 6.81 sec
MapReduce Total cumulative CPU time: 6 seconds 810 msec
Ended Job = job_1533488475608_0003
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 6.81 sec HDFS Read: 13992 HDFS Write: 263 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 810 msec
OK
Amitabh 101 2
Sharukh 102 1
Anubahv 104 1
Pawan 105 1
Aamir 106 1
Salman 107 1
Ranbir 108 1
Time taken: 67.893 seconds, Fetched: 7 row(s)
```

2. Create a new table called TRANSACTIONS_COUNT. This table should have 3 fields - custid, fname and count. (Again to be done in module 8)

Ans:

ScreenShot:

```
hive> CREATE TABLE TRANSACTIONS_COUNT(  
  > custid INT,  
  > fname STRING,  
  > count INT  
  > )  
  > row format delimited fields terminated by '\t';  
OK  
Time taken: 1.399 seconds
```

3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This must be done in module 9).

Ans:

CREATE VIEW trans_count_view AS

select t.custno,c.fname,count(txnno) from **TRANSACTIONS** t join **CUSTOMER** c on t.custno=c.custid **group by** t.custno,c.fname;

(Creating a view to store the result of transaction count. With the help of this view data would be feeded into newly created table).

Note: Data could have been directly inserted from the query itself but i have created a view to reduce the code complexity.

ScreenShot:

```
hive> CREATE VIEW trans_count_view AS  
  > select t.custno,c.fname,count(txnno) from TRANSACTIONS t join CUSTOMER c on t.custno=c.custid group by t.custno,c.fname;  
OK  
Time taken: 6.199 seconds
```

FROM trans_count_view

INSERT INTO TRANSACTIONS_COUNT **SELECT** *;

(Inserting into TRANSACTIONS_COUNT table for the view created.)

select * from TRANSACTIONS_COUNT;

(Displaying contents of TRANSACTIONS_COUNT table)

ScreenShot:

```
hive> FROM trans_count_view
> INSERT INTO TRANSACTIONS_COUNT SELECT *;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180806223434_d6e80ce5-9711-4b8b-9b10-d501b4204d76
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-08-06 22:34:51 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-06 22:34:56 Dump the side-table for tag: 1 with group count: 8 into file: file:/tmp/acadgild/b11b2db4-b649-4dd9-bbd2-240a646ed94e/hive_2018-08-06_22-34-34_004_2604854635531492418-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile01--.hashtable
2018-08-06 22:34:56 Uploaded 1 File to: file:/tmp/acadgild/b11b2db4-b649-4dd9-bbd2-240a646ed94e/hive_2018-08-06_22-34-34_004_2604854635531492418-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile01--.hashtable (469 bytes)
2018-08-06 22:34:56 End of local task; Time Taken: 5.075 sec.
Execution completed successfully
MapReduceLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533488475608_0004, Tracking URL = http://localhost:8088/proxy/application_1533488475608_0004/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533488475608_0004
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-06 22:35:21,522 Stage-2 map = 0%, reduce = 0%
2018-08-06 22:35:36,429 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.0 sec
2018-08-06 22:35:53,979 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 8.75 sec
MapReduce Total cumulative CPU time: 8 seconds 750 msec
Ended Job = job_1533488475608_0004
2018-08-06 22:34:56 End of local task; Time Taken: 5.075 sec.
Execution completed successfully
MapReduceLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533488475608_0004, Tracking URL = http://localhost:8088/proxy/application_1533488475608_0004/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533488475608_0004
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-06 22:35:21,522 Stage-2 map = 0%, reduce = 0%
2018-08-06 22:35:36,429 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.0 sec
2018-08-06 22:35:53,979 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 8.75 sec
MapReduce Total cumulative CPU time: 8 seconds 750 msec
Ended Job = job_1533488475608_0004
Loading data to table acadgild.db.transactions_count
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 8.75 sec HDFS Read: 14715 HDFS Write: 177 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 750 msec
OK
Time taken: 83.313 seconds
hive> select * from TRANSACTIONS_COUNT;
OK
101 Amitabh 2
102 Sharukh 1
104 Anubhav 1
105 Pawan 1
106 Aamir 1
107 Salman 1
108 Ranbir 1
Time taken: 0.526 seconds, Fetched: 7 row(s)
```

4. Now let's make the TRANSACTIONS_COUNT table Hbase compliant. In the sence, use Ser Des And Storate handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This must be done in module 10)

Ans:

CREATE TABLE TRANSACTIONS_HBase

```
(  
  custid INT,  
  fname STRING,  
  count INT  
)
```

STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'

WITH serdeproperties

("hbase.columns.mapping"=":key,details:name,details:txn_count")

tblproperties("hbase.table.name"="TRANSACTIONS");

(Creating a table TRANSACTIONS in HBase with *details* as column family along with a TRANSACTIONS_HBase table in HIVE. The **rowkey**, **name** & **txn_count** of TRANSACTIONS table in HBase are mapping to **custid**, **fname** & **count** columns of TRANSACTIONS_HBase table in HIVE)

ScreenShot:

NOTE: Before create table command in HIVE.

HBase does not consists of TRANSACTIONS table.

```
hbase(main):018:0> list  
TABLE  
bulktable  
clicks  
customer  
dept_tbl  
employee  
htest  
people  
tl  
8 row(s) in 0.0190 seconds  
=> ["bulktable", "clicks", "customer", "dept_tbl", "employee", "htest", "people", "tl"]
```

After the above create table command in HIVE:

ScreenShot:

```
hive> use acadgilddb;
OK
Time taken: 0.041 seconds
hive> show tables;
OK
customer
trans_count_view
transactions
transactions_count
Time taken: 0.133 seconds, Fetched: 4 row(s)
hive> create table TRANSACTIONS_HBase
> (
>   custid INT,
>   fname STRING,
>   count INT
> )
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> with serdeproperties ("hbase.columns.mapping"=":key,details:name,details:txn_count")
> tblproperties("hbase.table.name"="TRANSACTIONS");
OK
Time taken: 1.725 seconds
hive> show tables;
OK
customer
trans_count_view
transactions
transactions_count
transactions_hbase
Time taken: 0.134 seconds, Fetched: 5 row(s)
hive> desc transactions_hbase;
OK
custid      int
fname      string
count      int
Time taken: 0.225 seconds, Fetched: 3 row(s)
```

HBase :

```
hbase(main):019:0> list
TABLE
TRANSACTIONS
bulktable
clicks
customer
dept_tbl
employee
htest
people
t1
9 row(s) in 0.0170 seconds
=> ["TRANSACTIONS", "bulktable", "clicks", "customer", "dept_tbl", "employee", "htest", "people", "t1"]
hbase(main):020:0> desc "TRANSACTIONS"
Table TRANSACTIONS is ENABLED
TRANSACTIONS
COLUMN FAMILIES DESCRIPTION
{NAME => 'details', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TT
L => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
1 row(s) in 0.1970 seconds
```

NOTE: If HBase **TRANSACTIONS** table is disabled & dropped at this point the **TRANSACTIONS_HBase** table is HIVE would also automatically get dropped.

5. Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically (This must be done in module 10)

Ans:

Using the same view as in Step 3 above to insert the data in newly created TRANSACTION_HBASE table.

FROM trans_count_view

INSERT INTO TRANSACTIONS_HBASE SELECT *;

ScreenShot:

```
hive> FROM trans_count_view
> INSERT INTO TRANSACTIONS_HBase SELECT *;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180807120237_c7849a01-d799-42fb-9af0-ebc36e36286d
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-08-07 12:02:57 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-07 12:03:00 Dump the side-table for tag: 1 with group count: 8 into file: file:/tmp/acadgild/4a6313be-7ac2-427b-83d6-008e41adf559/hive_2018-08-07_12-02-37_126_3774401790868693631-1/-local-10002/HashTable-Stage-4/MapJoin-mapfile01--.hashtable
2018-08-07 12:03:01 Uploaded 1 File to: file:/tmp/acadgild/4a6313be-7ac2-427b-83d6-008e41adf559/hive_2018-08-07_12-02-37_126_3774401790868693631-1/-local-10002/HashTable-Stage-4/MapJoin-mapfile01--.hashtable (469 bytes)
2018-08-07 12:03:01 End of local task; Time Taken: 3.11 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1

set mapreduce.job.reduces=<number>
Starting Job = job_1533488475608_0013, Tracking URL = http://localhost:8088/proxy/application_1533488475608_0013/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533488475608_0013
Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 1
2018-08-07 12:03:24,384 Stage-4 map = 0%, reduce = 0%
2018-08-07 12:03:40,238 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 4.25 sec
2018-08-07 12:03:55,445 Stage-4 map = 100%, reduce = 76%, Cumulative CPU 9.31 sec
2018-08-07 12:03:57,879 Stage-4 map = 100%, reduce = 100%, Cumulative CPU 11.13 sec
MapReduce Total cumulative CPU time: 11 seconds 130 msec
Ended Job = job_1533488475608_0013
MapReduce Jobs Launched:
Stage-Stage-4: Map: 1 Reduce: 1 Cumulative CPU: 11.13 sec HDFS Read: 21293 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 130 msec
OK
Time taken: 82.933 seconds
hive> select * from trans_count_view;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180807120502_ic100c7d-aba7-4f3c-b838-ef2d927c84fb
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-08-07 12:05:20 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-07 12:05:24 Dump the side-table for tag: 1 with group count: 8 into file: file:/tmp/acadgild/4a6313be-7ac2-427b-83d6-008e41adf559/hive_2018-08-07_12-05-02_716_5586750326835459835-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile11--.hashtable
2018-08-07 12:05:24 Uploaded 1 File to: file:/tmp/acadgild/4a6313be-7ac2-427b-83d6-008e41adf559/hive_2018-08-07_12-05-02_716_5586750326835459835-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile11--.hashtable (469 bytes)
2018-08-07 12:05:24 End of local task; Time Taken: 4.005 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
```



```

2018-08-07 12:05:20 Starting to launch local task to process map join; maximum memory = 518979584
2018-08-07 12:05:24 Dump the side-table for tag: 1 with group count: 8 into file: file:/tmp/acadgild/4a6313be-
e_2018-08-07_12-05-02_716_5586750326835459835-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile11--.hashtable
2018-08-07 12:05:24 Uploaded 1 File to: file:/tmp/acadgild/4a6313be-7ac2-427b-83d6-008e41adf559/hive_2018-08-0
835-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile11--.hashtable (469 bytes)
2018-08-07 12:05:24 End of local task; Time Taken: 4.005 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533488475608_0014, Tracking URL = http://localhost:8088/proxy/application_1533488475608_0014/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533488475608_0014
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-08-07 12:05:40,388 Stage-2 map = 0%, reduce = 0%
2018-08-07 12:05:54,587 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.21 sec
2018-08-07 12:06:08,649 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 7.29 sec
MapReduce Total cumulative CPU time: 7 seconds 290 msec
Ended Job = job_1533488475608_0014
MapReduce Jobs Launched:
Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 7.29 sec HDFS Read: 13985 HDFS Write: 263 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 290 msec
OK
101 Amitabh 2
102 Sharukh 1
104 Anubahv 1
105 Pawan 1
106 Aamir 1
107 Salman 1
108 Ranbir 1
Time taken: 67.123 seconds, Fetched: 7 row(s)
hive>

```

Within HBase 7 rows inserted simultaneously:

```

hbase(main):021:0> scan "TRANSACTIONS"
ROW COLUMN+CELL
0 row(s) in 0.1070 seconds

hbase(main):022:0> scan "TRANSACTIONS"
ROW COLUMN+CELL
101 column=details:name, timestamp=1533623637153, value=Amitabh
101 column=details:txn_count, timestamp=1533623637153, value=2
102 column=details:name, timestamp=1533623637153, value=Sharukh
102 column=details:txn_count, timestamp=1533623637153, value=1
104 column=details:name, timestamp=1533623637153, value=Anubahv
104 column=details:txn_count, timestamp=1533623637153, value=1
105 column=details:name, timestamp=1533623637153, value=Pawan
105 column=details:txn_count, timestamp=1533623637153, value=1
106 column=details:name, timestamp=1533623637153, value=Aamir
106 column=details:txn_count, timestamp=1533623637153, value=1
107 column=details:name, timestamp=1533623637153, value=Salman
107 column=details:txn_count, timestamp=1533623637153, value=1
108 column=details:name, timestamp=1533623637153, value=Ranbir
108 column=details:txn_count, timestamp=1533623637153, value=1
7 row(s) in 0.2230 seconds

```

6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

Ans: JAVA API's attached as separate files.

Note: Program files are properly documented for a detailed description of each instruction used within the program.


ScreenShot:

The HBase table:

```
hbase(main):021:0> scan "TRANSACTIONS"
ROW
0 row(s) in 0.1070 seconds


hbase(main):022:0> scan "TRANSACTIONS"
ROW
101 column=details:name, timestamp=1533623637153, value=Amitabh
101 column=details:txn_count, timestamp=1533623637153, value=2
102 column=details:name, timestamp=1533623637153, value=Sharukh
102 column=details:txn_count, timestamp=1533623637153, value=1
104 column=details:name, timestamp=1533623637153, value=Anubahv
104 column=details:txn_count, timestamp=1533623637153, value=1
105 column=details:name, timestamp=1533623637153, value=Pawan
105 column=details:txn_count, timestamp=1533623637153, value=1
106 column=details:name, timestamp=1533623637153, value=Aamir
106 column=details:txn_count, timestamp=1533623637153, value=1
107 column=details:name, timestamp=1533623637153, value=Salman
107 column=details:txn_count, timestamp=1533623637153, value=1
108 column=details:name, timestamp=1533623637153, value=Ranbir
108 column=details:txn_count, timestamp=1533623637153, value=1
7 row(s) in 0.2230 seconds
```

For Access HBase:




```
1 import java.io.IOException;
2
3 public class accessHbasePk {
4     public static void main(String[] args) throws IOException, Exception {
5         //Instantiating Configuration class
6         Configuration config = HBaseConfiguration.create();
7
8         //Instantiating HTable class
9         @SuppressWarnings("deprecation")
10         HTable table = new HTable(config, "TRANSACTIONS");
11
12         //Instantiating Get class
13         Get g = new Get(Bytes.toBytes("108"));
14
15         //Reading the data
16         Result result = table.get(g);
17
18         //Reading values from Result class object
19         byte [] name = result.getValue(Bytes.toBytes("details"), Bytes.toBytes("name"));
20
21         byte [] txn = result.getValue(Bytes.toBytes("details"), Bytes.toBytes("txn_count"));
22
23         //Printing the values
24         String user = Bytes.toString(name);
25         String count = Bytes.toString(txn);
26
27         System.out.println("Customer Name: " + user + ", Number of transactions: " + count);
28     }
29 }
```

OutPut:




```
<terminated> accessHbasePk [Java Application] /usr/java/jdk1.8.0_151/bin/java (Aug 7, 2018, 2:14:06 PM)
log4j:WARN No appenders could be found for logger (org.apache.hadoop.security.Groups).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Customer Name: Ranbir, Number of transactions: 1
```

For Scan HBase:



```
1 import org.apache.hadoop.hbase.HBaseConfiguration;
2
3 public class scanTxnCount {
4
5     public static void main(String[] args) throws Exception {
6
7         //Instantiating Configuration class
8         HBaseConfiguration conf = new HBaseConfiguration();
9         //Instantiating HTable class
10        HTable htable = new HTable(conf, "TRANSACTIONS");
11        //Instantiating the Scan class
12        Scan scan = new Scan();
13        //Getting the scan result
14        ResultScanner scanner = htable.getScanner(scan);
15        Result r;
16        while ((r = scanner.next()) != null) {
17            //assign row values in variable userId
18            String userId = Bytes.toString(r.getRow());
19            //assign column username values in name
20            String name = Bytes.toString(r.getValue("details".getBytes(), "name".getBytes()));
21            //assign column txn count values in count
22            String count = Bytes.toString(r.getValue(Bytes.toBytes("details"), Bytes.toBytes("count")));
23            System.out.println("RowKey: " + userId + ", User Name: " + name + ", Count: " + count);
24        }
25        //closing the scanner
26        scanner.close();
27        htable.close();
28    }
29 }
```

OutPut:



```
<terminated> scanTxnCount [Java Application] /usr/java/jdk1.8.0_151/bin/java (Aug 7, 2018, 2:19:34 PM)
log4j:WARN No appenders could be found for logger (org.apache.hadoop.hbase.HBaseConfiguration).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
RowKey: 101, User Name: Amitabh, Count: 2
RowKey: 102, User Name: Sharukh, Count: 1
RowKey: 104, User Name: Anubhav, Count: 1
RowKey: 105, User Name: Pawan, Count: 1
RowKey: 106, User Name: Aamir, Count: 1
RowKey: 107, User Name: Salman, Count: 1
RowKey: 108, User Name: Ranbir, Count: 1
```

End
